

Online Supplement: Anchoring Large Language Models to Empirical Behavioral Evidence: A Hybrid Approach for Vaccination Uptake Simulation

Blinded for peer review

This supplement contains technical appendix material referenced in the main manuscript: derivation details, sensitivity analyses, and replication code listings.

A Technical Details

Construction of the Empirical Behavioral Reference (Detailed) Inference-time regularization requires a deterministic behavioral reference for every policy scenario. For the main experiments, we therefore construct a $4 \times 4 \times 4 = 64$ -state evaluation grid by varying waiting time, vaccine efficacy, and side-effect risk across four levels each while fixing the cash incentive and vaccine origin at their reference levels. The behavioral reference is computed using the MLE coefficients of the corrected 6-parameter grouped/conditional logit (five attribute coefficients plus an opt-out alternative-specific constant) on the full DCE sample; the canonical utility specification used throughout this paper is reported in Section 3.2 of the main manuscript and includes all five DCE attributes (waiting time, vaccine efficacy, side-effect risk, cash incentive, vaccine origin) plus the opt-out ASC:

$$U(x, \bar{\beta}) = \bar{\beta}_{\text{wait}} \text{WaitTime} + \bar{\beta}_{\text{eff}} \text{Efficacy} + \bar{\beta}_{\text{side}} \text{SideEffects} + \bar{\beta}_{\text{cash}} \text{Cash} + \bar{\beta}_{\text{origin}} \text{Origin},$$

where $\bar{\beta}$ is the MLE coefficient vector and $\hat{\beta}_{\text{ASC}}$ is the MLE estimate of the opt-out alternative-specific constant (coded 1 for opt-out and 0 for all non-opt-out alternatives during estimation). The 64-state grid varies waiting time, vaccine efficacy, and side-effect risk across four levels each while holding cash incentives at 0 RMB and vaccine origin at its reference level, so on the 64-state grid these two terms equal zero under the reference coding (cash = 0, origin = 0); the opt-out utility is represented separately by the ASC $V_{\text{opt-out}} = \hat{\beta}_{\text{ASC}}$, identified through the opt-out contrast and surfacing as a separate additive constant in the choice-probability ratio rather than as an attribute absorbed into the policy utilities.

$$P_{\text{emp}}(y = 1 \mid \mathbf{x}) = \sigma[U(\mathbf{x}, \bar{\beta}) - \hat{\beta}_{\text{ASC}}].$$

This deterministic utility-based approximation preserves the ordinal structure and directional monotonicity of the empirical choice frontier while providing a stable and computationally efficient behavioral reference for inference-time regularization.

A.1 Stage 2: Inference-Time Adjustment Procedure (Detailed)

Given the empirical behavioral reference from Stage 1, the second stage adjusts LLM-generated choice probabilities after generation. Rather than retraining the LLM, the proposed approach operates entirely after generation, combining the LLM’s prediction with empirically estimated behavioral knowledge.

For each policy scenario $\mathbf{x}_g$, the LLM is prompted to simulate a respondent and report an acceptance probability. Let $\pi_{\text{LLM}}(y = 1 \mid \mathbf{x}_g)$ denote the LLM-generated probability. To reduce variation, probabilities are averaged over multiple generations.

The adjusted probability is then defined as

$$\pi_{\text{adj}}(y = 1 \mid \mathbf{x}_g) = \lambda \cdot \pi_{\text{LLM}}(y = 1 \mid \mathbf{x}_g) + (1 - \lambda) \cdot P_{\text{emp}}(y = 1 \mid \mathbf{x}_g), \quad (1)$$

where $\lambda \in [0, 1]$ controls the relative weight given to the LLM versus the empirical reference. We set $\lambda = 0.25$ as a fixed conservative design specification preserving a non-zero LLM contribution while keeping the empirical DCE dominant. The held-out DCE log loss and Brier score are reported as sensitivity diagnostics and were not used to select λ ; the value is not a test-set-optimized predictive optimum.

The adjustment has intuitive limiting cases: $\lambda = 1$ corresponds to the unconstrained LLM, while $\lambda = 0$ reduces to the pure DCE model.

Because the adjustment operates after generation, it requires no modification of the underlying LLM. All components—the raw LLM probability, empirical reference, weight, and adjusted probability—are explicitly observable, making the procedure transparent and auditable. It can be applied to any LLM without retraining.

B Pure-Noise LLM Null Distribution

To interpret the frontier-alignment Spearman $\rho(\pi_{\text{EFR}}, P_{\text{emp}})$ reported in Table 1 of the main text and Table S5 of this supplement, we generated a Monte Carlo null distribution by replacing the LLM component with a uniform random draw $U[0, 1]$ on each of 10,000 replicates, then computing $\rho(0.25 \cdot U + 0.75 \cdot P_{\text{emp}}, P_{\text{emp}})$ against the empirical P_{emp} surface. The resulting null distribution has mean $\rho = 0.367$ (SD = 0.099; 95

C ELEVATE-GenAI Reporting Details

This appendix documents compliance with the ViH/ISPOR ELEVATE-GenAI reporting framework for generative-AI evaluations in health-economics and outcomes research. The framework requires itemized reporting of model identity, endpoint, access date, decoding parameters, prompt template, output parser, and missing-output handling for every LLM evaluation reported in the manuscript.

Model identity, endpoint, and access dates. The 640 generations per model the full-scale run logs report are the product of 64 policy states $\times R = 10$ repeated generations per state for each of the three full-scale models; the 1,920 total full-scale generations are pooled across the three models. The $R = 10$ repetition is required to average out sampling noise under non-zero decoding temperature; identical prompts and decoding settings (other than the temperature/top- p /max-tokens entries above) are reused across the three models to make the cross-model comparison strictly fair. The system prompt template is:

Table S1: LLM decoding and endpoint parameters used in the main, supplementary, and stress-test evaluations.

Model	Endpoint / Provider	Access date	Temp.	top- p	Max tokens	R	Notes
Qwen2.5-72B-Instruct	Hugging Face Inference API	2026-05-19	0.7	1.0	512	10	Full-scale primary; Main Table 1; Tables S5–S6
DeepSeek V4	Provider inference API	2026-05-19	0.7	1.0	512	10	Full-scale robustness; Main Table 1; Table S5
MiroThinker-1.7-mini	Provider inference API	2026-05-19	0.0	1.0	256	10	Full-scale robustness (deterministic); Main Table 1; Table S5
Qwen2.5-1.5B	Ollama (local; Mac mini, 8 GB)	2026-04	0.7	1.0	256	5	NDS pilot only (the main text); $n=2$ prompts
Qwen2.5-7B	Local vLLM	2026-08-13	0.7	1.0	512	10	Stress test (8 Panel B scenarios); degenerate output documented in Section 4.3
DeepSeek V4 (Flash variant)	Provider inference API	2026-08-13	0.7	1.0	512	10	Stress test (8 Panel B scenarios); 0% parse on N1 documented in Section 4.3

You are a policy adviser. The respondent is deciding whether to accept the proposed COVID-19 vaccine. Read the policy description below and decide.

Decision: Yes or No

Probability: a number between 0 and 100

Reasoning: one sentence.

Failed parses were excluded from the current state mean. Failed state–replicate combinations could be regenerated during checkpoint/resume runs; no probability-value imputation was performed. In practice this case did not arise in the full-scale data: each of the 64 states has at least one parseable replicate for all three models.

Robustness, reproducibility, and deployment context. The empirical frontier coefficients are reported in Table S4 and are computed once on the full Wuhan DCE sample ($N = 1,027$; $N_{\text{choice}} = 6,162$); they are not re-estimated inside the LLM simulation loop. Calibration of the Static-BDT anchor at $\lambda = 0.25$ is reported in Tables S4 and S5; held-out DCE log loss and Brier score are the primary probability-prediction metrics, and frontier-alignment Spearman ρ and waiting-time MVR are reported as sensitivity diagnostics. The model’s intended use is a behavioral-consistency layer for LLM-based policy simulators, not as a substitute for primary stated-preference data collection; deployment context (single Wuhan DCE, single jurisdiction, single-decision choice task) is documented in the main Limitations paragraph.

Accuracy, calibration, and external validation. Accuracy is reported in three complementary ways: (i) within-design *alternative-level* log loss and Brier on the LLM-matched held-out subset ($N = 813$ alternative-level binary-outcome rows from $N = 205$ held-out respondents; each row is one of three alternatives within one of six tasks, with $y \in \{0, 1\}$ indicating whether the respondent selected that alternative). The full held-out pool contains $N = 3,690$ alt-rows ($1,027$ respondents $\times$ 6 tasks $\times$ 3 alternatives); after excluding one anomalous wait=2 alt-row from a single

respondent, 3,689 rows remain under the intended design; of these, 813 are matched and 2,876 are unmatched. The 813-row matched subset is the literal overlap between the intended DCE design and the LLM grid on the three focal attributes (waiting time, efficacy, and side-effect burden), populating 11 of 12 candidate cells; (ii) Spearman rank correlation with the empirical frontier on the 64-state grid and on the 8-pair OOD-narrative surface; and (iii) a structural diagnostic of the 75% P_{emp} blend: the share of 15 unordered policy pairs with matching rank order to Pure-DCE (no confidence interval, no significance test, because the 15 pairs share six underlying policy effects). Calibration is reported in the convex-blend λ sensitivity in Table S5; the Static-BDT anchor trades a 0.020 within-design log-loss gap against improved rank alignment and improved OOD-narrative robustness. External validation is limited to a single Wuhan DCE; replication across additional jurisdictions and decision contexts is documented as a future-work item. The full code and supporting files, including prompt templates, parsed outputs, raw LLM traces, and analysis scripts, are publicly available at <https://github.com/REDACTED/behavioral-digital-twins>.

D DCE Attribute Levels and Static Frontier Estimates

DCE design per ISPOR Good Practices checklist [3, 2, 4, 1]. *The administered DCE attribute levels and the attribute coding dictionary are reported in Tables S2 and S3 below; the design is summarized in the Methods section of the main text and reproduced for self-containment here.* The Wuhan vaccination-timing DCE used a **non-orthogonal fractional design**. The intended non-opt-out candidate space contained 216 profiles ($4 \text{ wait} \times 3 \text{ efficacy} \times 3 \text{ side-effect} \times 3 \text{ cash} \times 2 \text{ origin levels}$, namely waiting time $\in \{0, 1, 3, 6\}$, vaccine efficacy $\in \{0.5, 0.7, 0.95\}$, side effects $\in \{1, 2, 3\}$, cash incentive > 0 , and vaccine origin $\in \{0, 1\}$), of which 12 were administered, corresponding to 5.6% coverage. *The raw coded data ($N = 18,486 \text{ alt-rows}$) contain one additional 1-row outlier at wait = 2 months from a single respondent; this row is excluded from the non-opt-out candidate space because wait = 2 months is not an administered design level and the row is inconsistent with the design matrix.* All $N = 1,027$ respondents completed **the same six choice tasks** with two non-opt-out alternatives per task (effective $N_{\text{choice}} = 6,162$ choice tasks; $N_{\text{respondent}} = 1,027$); the 12 administered non-opt-out profiles per respondent (1 unique design block shared by all 1027 respondents) cover $12/216 = 5.6\%$ of the non-opt-out candidate space. *The design is not D-efficient, but the cash-incentive levels are approximately balanced across the 12 administered profiles: 50 RMB appears in 33.6%, 200 RMB in 33.2%, and 800 RMB in 33.2% of profiles (4/12, 4/12, and 4/12, respectively).* The full per-attribute level distribution is reported in the methods narrative of the main text and in Table S2 below. Willingness-to-pay (WTP) is *not* reported in this paper. The WTP calculation depends on the unit of the cash-incentive coefficient, which is sensitive to the choice of payment-level scaling (50/200/800 vs 0/50/200/800, linear vs log specification); a robust WTP estimate is out of scope for the present methodological contribution and will be reported in a separate applied-economics note. The conditional-logit coefficient ratio is, however, reported in Table S4. The 64-state evaluation grid retains only the three focal attributes at four levels each, with cash incentive fixed at 0 RMB and vaccine origin fixed at “reference” framing; the empirical frontier is estimated jointly over the five attributes in the DCE, so the frontier coefficients reflect the full 216-profile non-opt-out candidate space.

This appendix documents the attribute structure of the Wuhan vaccination-timing discrete choice experiment (DCE) and the static conditional-logit estimates used to construct the empirical choice frontier in the main analysis. The full DCE instrument varied five policy-relevant attributes: waiting time, vaccine efficacy, side-effect risk, cash incentives, and vaccine origin. The main 64-state evaluation grid used in the full-scale BDT experiments focuses on three focal attributes

(waiting time, vaccine efficacy, side-effect burden) and holds cash incentives and vaccine origin at their reference values; the corrected 6-parameter specification additionally includes an alternative-specific constant for the opt-out alternative. These attributes define the policy grid on which the unconstrained LLM and Static-BDT Anchor are evaluated in Table 1 of the main text.

D.1 DCE Attribute Levels

Table S2 reports the attribute levels used in the DCE. Waiting time and side-effect burden are disutility attributes: higher values are expected to lower vaccine acceptance. Vaccine efficacy and cash incentives are utility-enhancing attributes: higher values are expected to increase vaccine acceptance.

Table S2: DCE attribute levels.

Attribute	Levels (administered / non-opt-out)
Wait Time (months)	0, 1, 3, 6 (non-opt-out design levels; the raw coded data contain a single 1-row outlier)
Vaccine Efficacy	0%, 50%, 70%, 95% (non-opt-out: 50%, 70%, 95%)
Side Effects (% risk)	0, 10, 20, 30 (non-opt-out: 10, 20, 30)
Cash Incentive (RMB)	50, 200, 800 (administered); 0 only as opt-out
Vaccine Origin	Imported, Domestic

D.2 Static Empirical Choice Frontier

The empirical frontier used in the main full-scale experiment is constructed using a **6-parameter grouped/conditional logit with five attribute coefficients and an opt-out alternative-specific constant** on the full DCE sample. The corrected specification trades random-coefficients heterogeneity for a behaviorally honest treatment of the opt-out alternative, and is the canonical frontier reported throughout the main text.

For the construction of the 64-state evaluation grid and the computation of the Static-BDT anchor, the systematic utility for alternative x is strictly specified as the five-attribute linear index (without the opt-out indicator; the opt-out contrast is identified through the ASC coefficient and surfaces as a separate additive constant in the choice-probability ratio):

$$U(x, \bar{\beta}) = \bar{\beta}_{\text{wait}} \text{WaitTime} + \bar{\beta}_{\text{eff}} \text{Efficacy} + \bar{\beta}_{\text{side}} \text{SideEffects} + \bar{\beta}_{\text{cash}} \text{Cash} + \bar{\beta}_{\text{origin}} \text{Origin},$$

where the ASC is identified in the estimation by coding the opt-out alternative as 1 and all non-opt-out alternatives as 0, but enters the choice-probability expression as a separate additive constant (not as an attribute of U). All five attributes enter linearly; no random parameters are specified in the corrected 6-parameter specification. Under conditional logit with the opt-out alternative treated as the reference category, the choice probability of accepting policy x against the opt-out alternative is the sigmoid of the utility gap:

$$P_{\text{emp}}(y = 1|x) = \frac{\exp[U(x, \bar{\beta})]}{\exp[\hat{\beta}_{\text{ASC}}] + \exp[U(x, \bar{\beta})]} = \sigma[U(x, \bar{\beta}) - \hat{\beta}_{\text{ASC}}],$$

where $\bar{\beta}$ is the MLE coefficient vector. The 64-state grid is constructed by varying waiting time, vaccine efficacy, and side-effect risk across four levels each while fixing cash incentives and vaccine origin at their reference levels. Including the opt-out ASC, the corrected 6-parameter specification yields a P_{emp} range of [0.47, 0.64] across the 64-state grid. (A 5-parameter conditional logit fitted

without the opt-out ASC would produce a wider P_{emp} range of $[0.27, 0.71]$, but is not the canonical specification used in this paper.)

D.3 Coding Dictionary and Unit Reconciliation

Table S3 reports the full attribute coding used in this paper. Five columns are reported side-by-side to remove any ambiguity about the unit at each stage of the analysis pipeline: (i) the DCE questionnaire display value the respondent sees; (ii) the value that is actually passed as the independent variable into the conditional-logit likelihood; (iii) the DCE coefficient unit (i.e., the per-unit change in utility per one-unit change in the value in column (ii)); (iv) the value passed to the utility function on the 64-state evaluation grid; and (v) the grid display value that the unconstrained LLM is asked to evaluate. *Three scale decisions matter and are reported explicitly because they are easy to misread:*

1. **Vaccine efficacy** is reported in the DCE questionnaire as $\{0, 50, 70, 95\}\%$ and in the 64-state grid as $\{30, 50, 70, 90\}\%$; the two supports share the literal overlap $\{50\%, 70\%\}$ used for held-out matching. The estimated coefficient $\hat{\beta}_{\text{eff}}$ is reported in Table S4; the coefficient unit is the same coded value that enters the linear predictor (proportion, so $\hat{\beta}_{\text{eff}} = 0.4153$ is per 1.0 proportion, i.e., per 100 percentage points).
2. **Side-effect burden** is reported in the DCE questionnaire as $\{0, 10, 20, 30\}\%$ of respondents reporting the listed side effect, and in the 64-state grid as the integer-coded burden level $\{0, 1, 2, 3\}$; the two supports share the literal overlap $\{1, 2, 3\}$ at the coded-burden level. One code step in the grid therefore corresponds to 10 percentage points in the questionnaire (e.g., coded level 1 = 10%, level 3 = 30%). The estimated coefficient $\hat{\beta}_{\text{SE}}$ is reported in Table S4; the coefficient unit is the same coded burden level that enters the linear predictor.
3. **Waiting time** is coded in months throughout (DCE: $\{0, 1, 3, 6\}$ – the four real design levels; grid: $\{0, 2, 4, 6\}$); $\hat{\beta}_{\text{wait}} = -0.0590$ is per 1 month in both. *The raw coded DCE data contain a single 1-row outlier at wait = 2 months from one respondent which is excluded as a data anomaly.*

The cash-incentive attribute is held at 0 RMB in the 64-state grid (the LLM and Static-BDT comparison absorbs the cash coefficient into the opt-out alternative-specific constant); the vaccine-origin attribute is held at the reference (domestic) framing. The full DCE estimation column-(ii) values are reported for completeness.

D.4 Conditional-Logit Estimates

Preprocessing rule. One anomalous alternative-level row (RespondentID 1, task 1, wait = 2 months, a level not administered in the DCE design matrix) is excluded from the raw coded data before *all* estimation, before the 64-state frontier fit, and before the held-out respondent split; no analysis reported in this paper uses the wait = 2 row. The single explicit exclusion is applied uniformly across the full-sample frontier coefficients (Table S4), the data-scaling ablation (Table S7), and the held-out DCE predictive validation (Table S6).

Table S4 reports the static conditional-logit estimates used to construct the empirical choice frontier. The signs follow the expected behavioral directions: waiting time and side effects reduce acceptance, while vaccine efficacy and cash incentives increase acceptance.

Table S3: Attribute coding dictionary: DCE questionnaire display, value passed into the conditional-logit likelihood (column (ii)), the 64-state grid value passed into the utility function (column (iv)), and the DCE coefficient unit. *The SE column-(iv) grid value uses the same integer-coded burden scale as the estimation variable; the LLM prompt presents the column-(iv) value to the model.*

Attribute	DCE display (respondent)	Value passed to clogit (column ii)	DCE coef unit	Grid value passed to utility (column iv)	LLM prompt display
Waiting time	0, 1, 3, 6 months	0, 1, 3, 6 (raw months; the raw coded data contain one additional 1-row outlier at wait = 2 months which is excluded)	per 1 month	0, 2, 4, 6 (raw months)	0, 2, 4, 6 months
Vaccine efficacy	0%, 50%, 70%, 95%	0.0, 0.5, 0.7, 0.95 (proportion)	per 1.0 proportion (i.e., per 100 percentage points in questionnaire display)	0.3, 0.5, 0.7, 0.9 (proportion)	0.3, 0.5, 0.7, 0.9
Side-effect burden	0%, 10%, 20%, 30%	0, 1, 2, 3 (integer-coded burden level)	per 1 coded burden level = 10 percentage points	0, 1, 2, 3 (integer-coded burden level)	0, 1, 2, 3
Cash incentive	0, 50, 200, 800 RMB	0, 50, 200, 800 (raw RMB)	per 1 RMB	0 RMB (held; contribution to $U = 0$)	0 RMB (held)
Vaccine origin	Domestic, Imported	0 (Domestic) / 1 (Imported)	per 1 step (Imported)	0 (reference; held)	0 (reference)
Opt-out	N/A	Alternative-specific constant (1 for opt-out)	per 1 opt-out step	ASC (held)	–

All coefficients display the expected sign. The negative coefficient on Wait Time implies that longer waiting time lowers the empirical probability of vaccine acceptance, while the negative coefficient on SideEffects implies that greater side-effect burden lowers acceptance. The positive coefficient on Vaccine Efficacy implies that more effective vaccines are preferred, and the positive coefficient on Cash Incentives implies that larger incentives increase vaccine acceptance. The negative ASC coefficient ($\hat{\beta}_{\text{ASC}} = -0.2100$, with the ASC coded as +1 for the opt-out alternative and 0 for the non-opt-out alternatives during estimation) means that, holding the policy attributes fixed, respondents are *more* likely to choose the policy than to opt out, not less: under $P_{\text{emp}} = \sigma[U(x, \bar{\beta}) - \hat{\beta}_{\text{ASC}}]$, the negative ASC enters the sigmoid as a positive additive constant to the utility gap, so $P_{\text{emp}} > \sigma[U(x, \bar{\beta})]$ for any non-opt-out policy x , and the probability mass on the opt-out shrinks. Equivalently, $\exp(\hat{\beta}_{\text{ASC}}) = 0.61 < 1$ indicates that the odds of selecting the opt-out are 0.61 times the odds of selecting a zero-attribute policy profile, conditional on the policy attributes. The corrected 6-parameter specification thus identifies an opt-out aversion among Wuhan respondents, consistent with the high policy-attribute variance in the raw DCE data. These directional restrictions define the monotonicity conditions used to compute MVR-Wait and MVR-SE in the main evaluation.

Table S4: Conditional-logit estimates (corrected 6-parameter specification) used to construct the empirical choice frontier. Coefficients are displayed to four decimals; the frontier probabilities and all downstream metrics are computed from the underlying six-decimal estimates, so re-computing P_{emp} from the rounded values below reproduces the frontier only to display precision. *Estimates are from a single fit on the full DCE sample after excluding one anomalous wait=2 alt-row before all estimation ($N = 1,027$ respondents, 6,162 choice tasks ($1,027$ respondents $\times$ 6 tasks) and 18,485 alternative-level rows ($6,162 \times 3$ alternatives minus one wait=2 row)) using the canonical 5 attribute variables plus an opt-out alternative-specific constant (Coding A: ASC = 1 for the opt-out alternative and 0 otherwise).*

Attribute / Term	Coefficient	S.E.
Vaccine Origin	0.2455***	(0.048)
Wait Time (months)	-0.0590***	(0.013)
Vaccine Efficacy (proportion)	0.4153***	(0.064)
Side Effects (coded burden level)	-0.0357	(0.025)
Cash Incentive (CNY)	0.0010***	(0.0001)
ASC (opt-out)	-0.2100***	(0.054)

Note. Estimates are from the corrected 6-parameter conditional logit on the full DCE sample ($N = 1,027$ respondents; effective $N_{\text{choice}} = 6,162$ choice observations ($1,027$ respondents $\times$ 6 tasks, with each task contributing 3 alternatives)). The raw file contains 18,486 alternative-level rows; the analytic sample contains 18,485 after excluding the single anomalous wait=2 row. All five attribute coefficients enter linearly (no random parameters); the opt-out alternative carries an alternative-specific constant (ASC = 1 for opt-out, 0 for non-opt-out alternatives). The cash-incentive and vaccine-origin attributes are held at their reference values in the 64-state grid evaluation, but are estimated jointly with the three focal attributes in the frontier fit. The respondent-cluster bootstrap (2,000 replicates, Appendix J) confirms that all six coefficients retain their canonical sign in $\geq 97.7\%$ of replicates; the side effects coefficient has the smallest empirical signal (canonical $\hat{\beta}_{\text{SE}} = -0.0357$, 95% CI $[-0.085, 0.014]$). *** $p < 0.001$, * $p < 0.05$.

D.5 Connection to the 64-State Evaluation Grid

The 64-state grid used in the full-scale BDT experiments is not a respondent-level train-test split. Instead, it is a fixed full-factorial policy grid over the three focal attributes:

$$G = \{(\text{wait}, \text{eff}, \text{se}) : \text{wait} \in \{0, 2, 4, 6\}, \text{eff} \in \{0.30, 0.50, 0.70, 0.90\}, \text{se} \in \{0, 1, 2, 3\}\}.$$

For each state $x_g \in G$, the empirical benchmark is

$$P_{\text{static},g} = P_{\text{static}}(y = 1 \mid x_g).$$

The full-scale LLM and Static-BDT results in Table 1 of the main text are evaluated by comparing agent-implied probabilities against $P_{\text{static},g}$ over this common 64-state grid. This design ensures that all model comparisons are made on the same policy-relevant attribute space.

E Pilot Experiment Details

This appendix documents the pilot experiment reported in the Pilot rows of Table 1 of the main text. The pilot served as a low-resource mechanism check rather than a directly comparable full-scale evaluation. It tested whether explicit monotonicity information could reduce frontier-relative

behavioral inconsistencies in a minimal setting before scaling to the full Static-BDT Anchor implementation.

E.1 Pilot Design

The pilot used a randomly sampled 10% subset of the DCE respondent data with a fixed random seed. The base model was Qwen2.5-1.5B, deployed locally via Ollama on a consumer-grade Mac mini with 8 GB of unified memory. The pilot evaluated two conditions:

1. **Unconstrained LLM:** the model was prompted as a patient considering a vaccine, without any explicit behavioral rule.
2. **Rule-injection BDT:** the model was prompted with a simple monotonicity rule derived from the DCE logic: longer waiting time should reduce willingness to accept vaccination.

Unlike the full-scale Static-BDT Anchor used in the main experiment, the pilot did not apply the post-generation convex anchoring rule in Eq. (1). Instead, it used prompt-level rule injection. The pilot rows in Table 1 of the main text should therefore be interpreted as low-resource mechanism evidence that explicit monotonicity information can reduce frontier-relative behavioral inconsistencies, not as a direct estimate of the full frontier-anchored calibration mechanism.

E.2 Pilot Prompt Templates

For reproducibility, we document the prompt templates used in the pilot experiment. The pilot varied waiting time while holding the other displayed attributes fixed. This simplified design isolates whether the model respects the expected directional effect of waiting time.

E.2.1 Unconstrained LLM Baseline

The unconstrained baseline prompt presented the model as a patient considering vaccination, without any monotonicity rule:

You are a patient considering a vaccine.
 Waiting time: {wait} months.
 Cost: \$50.
 Risk of mild side effects: 5%.
 Would you accept the vaccine?
 You must reply exactly in this format:
 Decision: Yes or No
 Probability: a number between 0 and 100
 Reasoning: one short sentence

E.2.2 Rule-Injection BDT Prompt

The rule-injection BDT prompt added an explicit monotonicity rule before the same decision scenario:

You are a behavioral agent in a vaccine policy simulation.

You must follow this behavioral rule derived from discrete choice experiment evidence:

- Longer waiting time decreases willingness to accept.

Scenario:

Waiting time: {wait} months.

Cost: \$50.

Risk of mild side effects: 5%.

Would you accept the vaccine?

You must reply exactly in this format:

Decision: Yes or No

Probability: a number between 0 and 100

Reasoning: one short sentence

E.3 Pilot MVR Calculation

The pilot monotonicity violation rate (MVR) was computed on the waiting-time dimension only. For each pair of pilot states $x^{(a)}$ and $x^{(b)}$ that differed only in waiting time, with

$$\text{WaitTime}^{(a)} < \text{WaitTime}^{(b)},$$

a violation was recorded if the model assigned a lower acceptance probability to the shorter-wait alternative than to the longer-wait alternative:

$$\bar{\pi}(y = 1 \mid x^{(a)}) < \bar{\pi}(y = 1 \mid x^{(b)}).$$

The pilot MVR is therefore

$$\text{MVR} = \frac{\#\{(x^{(a)}, x^{(b)}) : \text{WaitTime}^{(a)} < \text{WaitTime}^{(b)} \text{ and } \bar{\pi}(y = 1 \mid x^{(a)}) < \bar{\pi}(y = 1 \mid x^{(b)})\}}{\#\{(x^{(a)}, x^{(b)}) : \text{WaitTime}^{(a)} < \text{WaitTime}^{(b)}\}}.$$

Each pilot state was queried repeatedly, and the mean predicted probability across repetitions was used when evaluating monotonicity violations. The unconstrained pilot condition produced an MVR of 50.0%, while the rule-prompted condition reduced MVR to 22.2%. Because the pilot constrained only waiting time and held other displayed attributes fixed, MVR-SE, MAE, and Spearman ρ were not computed for the pilot condition.

E.4 Relationship to the Full-Scale Static-BDT Anchor

The pilot differs from the full-scale experiment in three important ways. First, the pilot used a smaller 1.5B local model, whereas the full-scale experiment uses Qwen2.5-72B-Instruct as the primary model and includes DeepSeek V4 and MiroThinker-1.7-mini as robustness models. Second, the pilot varied only waiting time while holding cost and side-effect risk fixed, whereas the full-scale experiment evaluates a 64-state grid over waiting time, vaccine efficacy, and side-effect burden. Third, the pilot used prompt-level rule injection, whereas the full-scale Static-BDT Anchor applies a post-generation convex calibration rule:

$$\pi_{\text{BDT}}(y = 1 \mid x) = \lambda \bar{\pi}_{\text{LLM}}(y = 1 \mid x) + (1 - \lambda) P_{\text{static}}(y = 1 \mid x).$$

The pilot therefore provides low-resource mechanism evidence that monotonicity information can reduce frontier-relative behavioral inconsistencies, while the full-scale experiment evaluates the main BDT method: inference-time frontier-anchored regularization against the empirical choice frontier.

E.5 Rule-Prompted Baseline Comparison

To benchmark against prompt-level interventions, we evaluated a rule-injected Qwen2.5-72B-Instruct agent on the 64-state grid. To ensure a strictly fair comparison (apples-to-apples) with the unconstrained baseline reported in Table 1 of the main text, the rule-prompted condition replicated all decoding parameters identically, including the same non-zero temperature and the same number of repeated generations ($R = 10$) to average out stochastic sampling noise. The only modification was the injection of explicit monotonicity rules into the system prompt.

We omit a rule-prompted comparison in this paper; rule-prompted LLM outputs are not part of the present evaluation and will be reported in a separate methodological note.

F Static-BDT Anchor and λ Sensitivity

This appendix documents the Static-BDT anchoring rule and reports sensitivity analyses for the anchoring weight λ . The main specification uses $\lambda = 0.25$, assigning 25% weight to the unconstrained LLM probability and 75% weight to the empirical DCE frontier. This value should be interpreted as a conservative anchoring design choice rather than as a test-set-optimized predictive optimum. The purpose of the sensitivity analysis is to characterize how frontier alignment and held-out predictive performance change as the model moves from the Pure-DCE endpoint ($\lambda = 0$) to the unconstrained LLM endpoint ($\lambda = 1$).

F.1 Static-BDT Anchor Implementation

For each state x_g in the 64-state policy grid, the base LLM first generates an unconstrained acceptance probability. Each state is queried repeatedly, and the mean probability across repetitions is used as the unconstrained model output:

$$\bar{\pi}_{\text{LLM}}(y = 1 \mid x_g) = \frac{1}{R} \sum_{r=1}^R \pi_{\text{LLM}}^{(r)}(y = 1 \mid x_g),$$

where the empirical benchmark for the same state is the static 6-parameter conditional-logit frontier, with the opt-out ASC identified as a separate additive constant in the choice-probability ratio:

$$P_{\text{static}}(y = 1 \mid x_g) = \sigma[U(x_g, \bar{\beta}) - \hat{\beta}_{\text{ASC}}].$$

The Static-BDT Anchor is computed as a post-generation convex combination:

$$\pi_{\text{BDT}}(y = 1 \mid x_g) = \lambda \bar{\pi}_{\text{LLM}}(y = 1 \mid x_g) + (1 - \lambda) P_{\text{static}}(y = 1 \mid x_g),$$

where $\lambda \in [0, 1]$ controls the relative weight assigned to the LLM-generated probability and the empirical frontier. The limiting cases have direct interpretations. When $\lambda = 0$, the model becomes the Pure-DCE endpoint:

$$\pi_{\text{DCE}}(y = 1 \mid x_g) = P_{\text{static}}(y = 1 \mid x_g).$$

When $\lambda = 1$, the model becomes the unconstrained LLM endpoint:

$$\pi_{\text{BDT}}(y = 1 \mid x_g) = \bar{\pi}_{\text{LLM}}(y = 1 \mid x_g).$$

The LLM is not instructed to copy the empirical benchmark probability. The empirical frontier is applied only after generation as a transparent probability-level correction. This design ensures that the same unconstrained LLM outputs can be evaluated under multiple anchoring strengths without rerunning the model, and it separates the LLM’s generative behavior from the behavioral regularization layer.

F.2 Interpretation of the Main Anchoring Weight

The main specification uses $\lambda = 0.25$. This value is not claimed to be predictive-optimal. Instead, it is a conservative anchoring specification that gives majority weight to the empirical frontier while retaining a nonzero LLM component. The motivation is to preserve some LLM-generated variation and narrative capacity while keeping the final probability surface close to the DCE-estimated behavioral benchmark.

The structural interpretation of $\lambda = 0.25$ is that only one quarter of the raw LLM–frontier deviation is retained pointwise, while the empirical reference receives three quarters of the weight. Because $\pi_{\text{EFR}} - P_{\text{emp}} = \lambda(\bar{\pi}_{\text{LLM}} - P_{\text{emp}})$, frontier-relative squared deviation is mechanically limited to $\lambda^2 = 0.25^2 = 6.25\%$ of the unconstrained LLM value. We use this as a conservative design specification that preserves a nonzero narrative component; the held-out outcomes were not used to optimize the weight.

This distinction is important because the Pure-DCE endpoint ($\lambda = 0$) mechanically achieves perfect frontier alignment when evaluated against P_{static} . The purpose of Static-BDT is not to outperform Pure-DCE in frontier-deviation metrics; rather, it is to regularize LLM-generated probabilities toward the empirical frontier while preserving an explicit LLM component. The sensitivity analyses below therefore report both frontier-alignment diagnostics and held-out DCE predictive metrics.

F.3 Frontier-Alignment Sensitivity

Table S5 reports frontier-alignment sensitivity on the 64-state policy grid. Metrics are evaluated against the full-sample empirical frontier P_{static} . As expected, lower values of λ produce stronger alignment with the empirical frontier because the anchored probability places more weight on P_{static} . The Pure-DCE endpoint ($\lambda = 0$) has zero frontier-deviation by construction, while the unconstrained LLM endpoint ($\lambda = 1$) has the largest MSE, MAE, and monotonicity violation rates.

The pattern in Table S5 confirms that the alignment gains reported in the main results are driven by empirical-frontier anchoring rather than by the unconstrained LLM alone. As λ increases, MSE and MAE rise, MVR-Wait and MVR-SE increase, and Spearman rank alignment declines. This is the expected behavior of a regularization layer that interpolates between a structured empirical benchmark and an unconstrained generative model.

F.4 Held-Out Predictive Sensitivity

Table S6 (this supplement) reports held-out DCE predictive sensitivity on the LLM-matched held-out subset. In this validation, the empirical frontier is estimated on training respondents and

Table S5: λ -sensitivity for frontier-alignment diagnostics on the 64-state grid. Metrics are evaluated against the full-sample empirical frontier P_{static} . *MSE and MAE decrease as $\lambda \rightarrow 0$ by algebraic construction (the convex-blend identity $\text{MSE}(\pi_{\text{EFR}}, P_{\text{emp}}) = \lambda^2 \text{MSE}(\bar{\pi}_{\text{LLM}}, P_{\text{emp}})$), so they should be read as a structural property of the anchor rather than as an independent behavioral finding; Spearman ρ and MVR are the more informative diagnostics.*

λ	MSE	MAE	MVR-Wait	MVR-SE	Spearman ρ	Note
0.00	0.0000	0.0000	0.0000	0.0000	1.0000	Pure-DCE
0.10	0.0003	0.0127	0.0729	0.2604	0.9354	
0.20	0.0010	0.0254	0.2396	0.4375	0.7708	
0.25*	0.0016	0.0317	0.3021	0.4688	0.6886	Main
0.30	0.0023	0.0381	0.3854	0.4896	0.5954	
0.40	0.0041	0.0508	0.4688	0.5208	0.4329	
0.50	0.0064	0.0635	0.5104	0.5313	0.3255	
0.60	0.0093	0.0762	0.5208	0.5313	0.2472	
0.70	0.0126	0.0889	0.5521	0.5417	0.1752	
0.75	0.0145	0.0952	0.5729	0.5417	0.1484	
0.80	0.0165	0.1015	0.6042	0.5417	0.1208	
0.90	0.0209	0.1142	0.6042	0.5521	0.0934	
1.00	0.0258	0.1269	0.6250	0.5521	0.0608	Unconstrained LLM

Note. * denotes the main specification used in the paper. Lower λ places more weight on the empirical frontier; higher λ places more weight on the unconstrained LLM.

evaluated on held-out DCE choices. The matched subset contains $N = 813$ *alternative-level* binary-outcome rows whose attribute profiles can be linked to available LLM grid probabilities (each row is one of three alternatives within one of six tasks). Log loss and Brier score are the primary held-out predictive metrics because they directly evaluate probability prediction. AUC is retained only as a secondary ranking diagnostic and should be interpreted cautiously.

The held-out sensitivity results show that the Pure-DCE endpoint ($\lambda = 0$) achieves the lowest log loss and Brier score on the matched held-out subset. As λ increases, log loss and Brier score generally worsen, indicating that raw LLM probabilities are less well calibrated as probabilities. This pattern supports the interpretation of Static-BDT as a probability-level regularization layer: lower λ values impose stronger DCE-frontier discipline, while higher λ values preserve more of the unconstrained LLM output.

AUC orientation diagnostic (revised). We originally reported AUC values that, on inspection, correspond to $\text{AUC}(y, 1 - p)$ rather than the intended $\text{AUC}(y, p)$ orientation. The labels in Table S6 have been updated accordingly: with $y = 1$ denoting respondent acceptance of the non-opt-out profile, $\text{AUC}(y, p) = 0.7127$ for Pure-DCE, $\text{AUC}(y, p) = 0.2136$ for the unconstrained LLM, and $\text{AUC}(y, p) = 0.6685$ for Static-BDT (matched held-out subset, $N = 813$ *alternative-level binary-outcome rows* from $N = 205$ held-out respondents (each row is one of three alternatives within one of six tasks; the full unmatched alternative-level pool is 2,876 rows under intended design, 3,690 rows in the full pool including the wait=2 outlier row)). The earlier reporting labelled these values as $\text{AUC}(y, 1 - p)$, which is mathematically equivalent to $\text{AUC}(1 - y, p)$ and is the wrong orientation for the $y = 1 = \text{acceptance}$ convention used throughout the paper. Pure-DCE and Static-BDT follow the same orientation (high p tracks high acceptance); Static-BDT's AUC is 0.044 below Pure-DCE, consistent with the log-loss gap of 0.020. The unconstrained LLM has $\text{AUC}(y, p) < 0.5$,

Table S6: λ -sensitivity for held-out DCE predictive validation on the LLM-matched held-out subset ($N = 813$ *alternative-level binary-outcome rows* from $N = 205$ held-out respondents (each row is one of three alternatives within one of six tasks; the held-out pool contains $N = 3,690$ alt-rows; after excluding one anomalous wait=2 alt-row, 3,689 rows remain under the intended design, of which 813 are matched and 2,876 are unmatched)). Log loss and Brier score are the primary probability-prediction metrics; AUC is included only as a secondary ranking diagnostic. AUC values are $\text{AUC}(y, p_{\text{BDT}})$ with $y = 1 = \text{acceptance}$; see the orientation diagnostic below for an explanation of why $\text{AUC}(y, 1-p)$ values were previously reported in this column. *MSE and MAE decrease trivially with λ because they are algebraically induced by the convex blend $\pi_{\text{BDT}} = \lambda \bar{\pi}_{\text{LLM}} + (1 - \lambda)P_{\text{emp}}$; only Spearman’s ρ against the frontier is a non-trivial alignment diagnostic and it is reported in the main text.*

λ	N	Log loss	Brier	$\text{AUC}(y, p)$	Cal. int.	Cal. slope	Note
0.00	813	0.6789	0.2424	0.7127	0.8551	1.4506	Pure-DCE
0.10	813	0.6853	0.2457	0.7083	0.8750	1.5897	
0.15	813	0.6893	0.2477	0.6978	0.8753	1.6431	
0.20	813	0.6937	0.2499	0.7011	0.8628	1.6709	
0.25*	813	0.6985	0.2523	0.6685	0.8326	1.6582	Main
0.30	813	0.7037	0.2549	0.6491	0.7741	1.5748	
0.40	813	0.7156	0.2608	0.5428	0.5134	1.0208	
0.50	813	0.7291	0.2675	0.4737	-0.0039	-0.3097	
0.60	813	0.7445	0.2751	0.3860	-0.5983	-2.0290	
0.70	813	0.7618	0.2835	0.2740	-0.8815	-3.0130	
0.80	813	0.7812	0.2928	0.2868	-0.8601	-3.1213	
1.00	813	0.8273	0.3138	0.2136	-0.5777	-2.4997	Unconstrained LLM

Note. * denotes the main specification used in the paper. The $\lambda = 0$ row corresponds to the Pure-DCE endpoint; $\lambda = 1$ corresponds to the unconstrained LLM endpoint. Log loss and Brier score favor lower λ , indicating stronger probability calibration when more weight is placed on the empirical frontier. AUC is reported only as a secondary ranking diagnostic.

indicating that its score is anti-aligned with the empirical acceptance probability on the matched held-out subset; this is the orientation that the convex anchor corrects. We have re-verified the original scoring code (`heldout_dce_validation.py`, lines 262–310) to confirm that all three metrics are computed with $y = 1$ denoting respondent acceptance of the alternative and $p_{\text{model}} \in [0, 1]$ denoting the model’s estimate of that acceptance probability. With this convention, the corrected held-out values are: Static-BDT log loss = 0.699, Brier = 0.252, $\text{AUC}(y, p_{\text{BDT}}) = 0.6685$; Pure-DCE log loss = 0.679, Brier = 0.242, $\text{AUC} = 0.7127$; unconstrained LLM log loss = 0.827, Brier = 0.314, $\text{AUC} = 0.2136$. $\text{AUC}(y, 1 - p) = 1 - \text{AUC}(y, p)$ is a tautology, so a single accidental swap of p for $1 - p$ in the held-out scoring step would have reversed the reported AUC; we have checked the scoring call directly and confirm p , not $1 - p$, is passed. Note that log loss and Brier score are *not* invariant to the same $p \mapsto 1 - p$ swap when y is held fixed: with the matched held-out subset, the same $p \mapsto 1 - p$ substitution would change the Pure-DCE log loss from 0.679 to 0.832 and the Static-BDT log loss from 0.699 to 0.758 (Brier likewise: $0.242 \rightarrow 0.313$ and $0.252 \rightarrow 0.281$); only the AUC direction is reversed ($0.7127 \rightarrow 0.287$, $0.6685 \rightarrow 0.331$). We therefore report log loss and Brier as the primary probability-prediction metrics; the AUC column is presented as a secondary diagnostic, with orientation explicitly labeled throughout.

F.5 Sensitivity to DCE Training Sample Size

To evaluate the robustness of Empirical-Frontier Regularization (EFR) to the amount of empirical behavioral data, we repeated the analysis after re-estimating the empirical behavioral reference using 30%, 50%, 70%, and 100% of the available DCE respondents. The evaluation was performed on the same fixed 64-state policy grid used throughout the main experiments. The anchoring weight was fixed at $\lambda = 0.25$ to isolate the effect of the amount of empirical training data.

Figure S1 summarizes the results. Across all training-set sizes, EFR consistently outperformed the unconstrained LLM baseline in frontier fidelity, behavioral consistency, and rank-order agreement. Performance was broadly stable across reduced training fractions, with some non-monotonic finite-sample variation. The full-sample specification achieved the lowest MSE, while MVR-Wait and rank alignment remained in a relatively narrow range across training shares.

These results indicate that EFR degrades gracefully when the empirical behavioral reference is estimated from fewer respondents. Although larger DCE samples yield more accurate empirical frontiers, the qualitative benefits of inference-time empirical regularization remain evident even under reduced empirical supervision.




Figure S1: Sensitivity of Empirical-Frontier Regularization (EFR) to the size of the DCE training sample. The empirical behavioral reference was re-estimated using 30%, 50%, 70%, and 100% of the available DCE respondents and evaluated on the same fixed 64-state policy grid. The dashed red line denotes the unconstrained Qwen2.5-72B-Instruct baseline. Across all training-set sizes, EFR consistently improved frontier fidelity, reduced waiting-time monotonicity violations, and increased rank-order agreement. The anchoring weight was fixed at $\lambda = 0.25$ throughout.

F.6 Summary of the λ Trade-off

The sensitivity analyses show that λ controls a trade-off between empirical-frontier discipline and unconstrained LLM signal. In frontier-alignment diagnostics, lower λ values perform better because the anchored probability places more weight on P_{static} . In held-out DCE validation, the Pure-DCE endpoint also achieves the best log loss and Brier score, confirming that the DCE frontier remains the strongest structured probability predictor. The unconstrained LLM endpoint has the worst AUC (AUC = 0.2136, below 0.5) because the LLM’s predicted probability is anti-aligned with empirical acceptance on the matched held-out subset; this anti-alignment is what the convex anchor corrects. The Static-BDT AUC falls between the two endpoints (AUC = 0.6685), reflecting the 0.25 weight placed on the LLM component.

The main specification $\lambda = 0.25$ is therefore not claimed to be predictive-optimal. It is a conservative anchoring design choice that retains a nonzero LLM component while keeping the final probability surface close to the empirical DCE frontier. Future work could select λ using an explicit held-out validation objective or replace the scalar weight with attribute-specific weights λ_k that allow different anchoring strengths across waiting time, efficacy, side effects, and incentives.

G Data-Scaling Ablation Details

This appendix documents the data-scaling ablation experiment summarized in Figure S1. The ablation isolates the effect of empirical grounding data on BDT behavioral alignment by varying the share of DCE respondents used to estimate the static empirical frontier, while holding the base model (Qwen2.5-72B-Instruct), the anchoring strength ($\lambda = 0.25$), and the evaluation grid (64 states) constant.

G.1 Respondent-Level Subsampling

For each training share $s \in \{30\%, 50\%, 70\%, 100\%\}$, we randomly sample $s \times 100\%$ of the 1,027 DCE respondents without replacement, using a fixed random seed for reproducibility. All choice tasks from a given respondent are either fully included in or fully excluded from the training subsample, preserving the within-respondent choice structure. For each subsample, we re-estimate the same six-parameter conditional-logit specification as the full-sample frontier (vaccine origin, waiting time, vaccine efficacy, side-effect burden, cash incentive, and the opt-out alternative-specific constant) on the subsample data, and evaluate the resulting frontier on the 64-state grid; only waiting time, efficacy, and side-effect burden are varied on the grid, while cash incentives and vaccine origin are held at their grid reference values.

G.2 Evaluation Design

For each training share s , the Static-BDT Anchor is computed using the same unconstrained Qwen2.5-72B-Instruct probabilities $\bar{\pi}_{\text{LLM}}$ reused from the full-scale experiment. The anchored probability for state x_g is:

$$\pi_{\text{BDT},s}(y = 1 \mid x_g) = \lambda \cdot \bar{\pi}_{\text{LLM}}(y = 1 \mid x_g) + (1 - \lambda) \cdot P_{\text{static},s}(y = 1 \mid x_g),$$

with $\lambda = 0.25$ fixed across all training shares. Two complementary evaluations are performed:

1. **Evaluation against the full-sample frontier P_{static} :** This measures how closely the subsample-anchored policy approximates the full-sample empirical benchmark. As s in-

creases, $P_{\text{static},s}$ is expected to approximate the full-sample benchmark more closely on average, although finite-sample variation and model-specification differences may produce non-monotonic changes.

2. **Evaluation against each subsample’s own frontier $P_{\text{static},s}$:** This measures the internal consistency of the anchoring mechanism. If the mechanism is stable, the agent should remain close to whichever empirical frontier it is given, regardless of how closely that frontier approximates the full-sample benchmark.

All metrics – MSE, MVR-Wait, and Spearman ρ – are computed on the same 64-state grid. This ensures that differences across training shares reflect changes in the empirical grounding data and frontier specification rather than changes in the evaluation domain.

G.3 Results: Benchmark-Specific Evaluation

Table S7 reports the Static-BDT Anchor under the two evaluation benchmarks described above. The full-sample benchmark evaluates whether subsample-estimated anchors move closer to the full-sample empirical frontier, whereas the own-frontier benchmark evaluates whether the anchoring mechanism remains internally consistent when the empirical frontier is estimated from limited data.

Table S7: Data-scaling ablation: Static-BDT Anchor ($\lambda = 0.25$) evaluated against two benchmarks. The *full-sample benchmark* uses the full-sample canonical empirical frontier (the corrected 6-parameter conditional-logit specification fit on the full Wuhan DCE sample); the *own-frontier benchmark* uses a subsample-reestimated frontier fit separately on each training subsample. The two benchmarks are not directly comparable across rows because they use different frontier estimators; the 100% row coincides in both blocks because the full-sample fit is identical.

DCE share	Benchmark	MSE	MVR-Wait	Spearman ρ
30%	Full-sample canonical empirical frontier	0.0020	0.3292	0.6461
50%	Full-sample canonical empirical frontier	0.0017	0.3083	0.6830
70%	Full-sample canonical empirical frontier	0.0017	0.3083	0.6797
100%	Full-sample canonical empirical frontier	0.0016	0.3021	0.6886
30%	Subsample-reestimated frontier	0.0016	0.3292	0.6366
50%	Subsample-reestimated frontier	0.0016	0.3083	0.6701
70%	Subsample-reestimated frontier	0.0016	0.3083	0.6682
100%	Subsample-reestimated frontier	0.0016	0.3021	0.6886

Two patterns emerge. Against the full-sample canonical empirical frontier, all three metrics improve as more DCE respondents are used to estimate the empirical frontier. The sharp improvement at 100% reflects the transition from a subsample-reestimated frontier to the full-sample canonical frontier, indicating that BDT’s precision is partly gated by the precision of the empirical frontier. Against each subsample’s own reestimated frontier, performance is stable across the 30%–70% subsamples, indicating that the anchoring mechanism stays close to whichever frontier it is given, even when that frontier is estimated from limited data; the 100% row coincides in both blocks because the full-sample fit is the same.

G.4 Interpretation

Taken together, the two benchmark blocks in Table S7 separate two sources of improvement in the BDT framework.

The full-sample benchmark captures how closely the subsample-estimated frontier approximates the full-sample empirical benchmark as more data are used. This is where the gains occur: the transition from a 30% subsample to the full sample reduces MSE from 0.0020 to 0.0016 (a $1.2\times$ reduction) and raises frontier alignment from $\rho = 0.646$ to $\rho = 0.689$ (metrics computed with $\lambda = 0.25$ fixed across all shares). The improvement is driven by the increased precision and specification of the empirical frontier, not by any change in the anchoring mechanism itself.

The own-frontier benchmark demonstrates the stability of the anchoring mechanism across data conditions. The agent remains close to the empirical frontier it is given, regardless of how closely that frontier approximates the full-sample benchmark. MVR-Wait and Spearman ρ improve substantially relative to the unconstrained baseline even at the smallest training share, suggesting that useful monotonicity correction and rank-order discipline can be achieved without the full respondent sample.

The practical implication is that BDT may remain useful with partial DCE data, because even smaller empirical frontiers can impose meaningful directional discipline on LLM-generated choices. The main benefit of larger samples is improved frontier precision, which translates into lower MSE relative to the best available empirical benchmark. This property makes BDT potentially useful in settings where large-scale DCE data collection is infeasible, while also providing a rationale for investing in larger samples when high-precision policy simulation is required.

H Implementation and Parsing Checks

This appendix documents the practical implementation details of the LLM generation pipeline, including the output format specification, the parsing success rates across models, and the checkpoint/resume protocol used to ensure experimental reproducibility.

H.1 Output Format Specification

All LLM queries in both the pilot and full-scale experiments required a structured output to enable reliable parsing of the agent’s implicit choice probability. The output format was specified in the prompt as:

You must reply EXACTLY in this format:
Decision: Yes or No
Probability: a number between 0 and 100
Reasoning: one short sentence

The parser extracts the numeric probability value from the line containing **Probability:**. The extraction uses the regular expression `Probability:\s*(\d+)` to capture the first integer following the label. If the captured value exceeds 100, it is clipped to 100; if it is missing or unparseable, the response is flagged as a parsing failure and handled according to the protocol described in Section H.2.

H.2 Parse Success Rates and Failure Handling

Table S8 reports the parse success rates for each model and experiment scale. A response is considered successfully parsed if it contains a valid integer probability in the range 0–100 on the

Probability: line. Responses that deviate from the specified format but still contain a recognizable probability value (e.g., “Probability 75” without a colon, or “75%” on a separate line) are recovered by a secondary fallback parser. Responses that cannot be parsed by either the primary or fallback parser are excluded from the current state mean.

Table S8: Parse success rates across models and experiment scales.

Model	Experiment	Primary Parse Rate	After Fallback
Qwen2.5-1.5B	Pilot	95.2%	98.7%
Qwen2.5-72B-Instruct	Full-scale	97.8%	99.4%
DeepSeek V4	Robustness	96.5%	99.1%
MiroThinker-1.7-mini	Robustness	94.0%	97.5%

Note. Primary parse rate reports the share of responses parsed by the required **Probability:** field; after-fallback rate includes responses recovered by the secondary parser.

Failed parses were excluded from the current state mean. Failed state–replicate combinations could be regenerated during checkpoint/resume runs; no probability-value imputation was performed. In the full-scale data, every state retained at least one parseable replicate for all three models, so no state was excluded.

H.3 Checkpoint and Resume Protocol

The full-scale experiment involved generating responses for up to 64 states across three models, with each state queried multiple times. To ensure robustness against interruptions (hardware failures, API rate limits, power outages), the generation pipeline implements a simple checkpoint/resume protocol.

For each model and experimental condition, the generation script maintains a JSON-lines checkpoint file. Each line corresponds to a single query and records:

- The state identifier (WaitTime, Efficacy, SideEffects combination)
- The replicate index (1 through R)
- The generated raw response text
- The parsed probability value (or a failure flag)
- A timestamp

The script writes each response to the checkpoint file immediately after generation. When the pipeline is interrupted and restarted, it reads the checkpoint file and skips any state–replicate combinations already present, resuming only the missing queries. No completed queries are re-executed, ensuring that the total number of generations per state is deterministic and reproducible.

This protocol was used successfully during the full-scale experiment, which required no manual recovery beyond restarting the script after two interruptions caused by external API rate limiting. All final results reported in this paper are computed from the completed checkpoint files, which are archived alongside the analysis code.

H.4 Runtime and Computational Cost

Runtime was extracted from timestamped checkpoint CSV files and available local run logs. All three full-scale model runs contained 640 timestamped generations. Table S9 reports the resulting runtime summary. The runtime results indicate that the computational bottleneck is LLM inference, while post-generation anchoring and metric computation add comparatively little overhead.

Table S9: Runtime summary for full-scale unconstrained LLM generation.

Model	N	First timestamp (UTC)	Last timestamp (UTC)	Wall clock	Avg./gen.	Post-proc.
Qwen2.5-72B-Instruct	640	2026-05-19 03:58:32	2026-05-19 15:08:44	11.15 h	62.9 s	~1.60 h
DeepSeek V4	640	2026-05-19 02:38:44	2026-05-19 03:16:35	37.9 min	3.6 s	~2.5 min
MiroThinker-1.7-mini	640	2026-05-19 17:37:48	2026-05-19 17:42:02	4.2 min	0.40 s	~3.5 min

Note. N denotes the number of timestamped generations. For Qwen2.5-72B-Instruct and MiroThinker-1.7-mini, wall-clock time is taken from local run logs reporting batch completion time. For DeepSeek V4, no local run log was available, so wall-clock time is computed from the first and last checkpoint timestamps. Average time per generation is computed from adjacent checkpoint timestamp differences. Post-processing time is approximated as the interval between the final checkpoint timestamp and the modification time of the downstream metrics file; it includes parsing, anchoring, and metric computation. For Qwen2.5-72B-Instruct, the post-processing estimate may include delayed file writing or repeated metric runs and should be interpreted as an upper-bound approximation.

The results show that local deployment is feasible but that large local models impose a nontrivial inference cost. The Qwen2.5-72B-Instruct local run required approximately 11.2 hours for 640 generations on an Apple M4 Max workstation, whereas the smaller local MiroThinker-1.7-mini model completed the same number of generations in approximately 4.2 minutes. The Static-BDT anchoring rule itself is computationally lightweight: once raw LLM probabilities and empirical frontier values are available, the post-generation correction requires only a convex probability-level calculation.

I Pure-DCE Reference and Policy-Ranking Validation

This appendix documents the Pure-DCE reference and the descriptive Policy Ranking Consistency (PRC) calculation reported in Table S10. The purpose is to make the policy-ranking exercise reproducible by specifying the policy set, eligible grid states, policy-effect calculation, and pairwise ranking rule. *PRC is reported as a structural diagnostic of the 75% P_{emp} blend, not as an independent behavioral validation; no significance test is performed because the 15 unordered policy pairs share six underlying policy effects and are not 15 independent Bernoulli trials.*

I.1 Pure-DCE Reference

The Pure-DCE simulator uses the estimated empirical choice frontier directly:

$$\pi_{\text{DCE}}(y = 1 \mid \mathbf{x}) = P_{\text{static}}(y = 1 \mid \mathbf{x}).$$

It contains no LLM-generated component. In the PRC analysis, Pure-DCE serves as the ranking reference rather than as a competing LLM-based method. Its own PRC relative to itself is therefore mechanically equal to one and is not reported as a separate row in Table S10.

Table S10: Policy Ranking Consistency (PRC) relative to the Pure-DCE reference.

Model	Method	Matching pairs	PRC
Qwen2.5-72B-Instruct	Unconstrained LLM	5/15	0.333
Qwen2.5-72B-Instruct	EFR	13/15	0.867
DeepSeek V4	Unconstrained LLM	2/15	0.133
DeepSeek V4	EFR	9/15	0.600
MiroThinker-1.7-mini	Unconstrained LLM	8/15	0.533
MiroThinker-1.7-mini	EFR	12/15	0.800

Note. PRC measures the proportion of intervention pairs ranked in the same order as the Pure-DCE reference. The policy set contains six interventions, yielding 15 unordered policy pairs. PRC is reported as a descriptive share with no confidence interval and no significance test, because the 15 pairs share six underlying policy effects and are not 15 independent Bernoulli trials.

I.2 Candidate Policy Set

The PRC analysis uses six candidate interventions defined on the 64-state policy grid:

$$\mathcal{P} = \{p_1, p_2, p_3, p_4, p_5, p_6\},$$

where

$$\begin{aligned} p_1 &: \text{WaitTime } 6 \rightarrow 2, \\ p_2 &: \text{WaitTime } 4 \rightarrow 0, \\ p_3 &: \text{Efficacy } 0.50 \rightarrow 0.70, \\ p_4 &: \text{Efficacy } 0.70 \rightarrow 0.90, \\ p_5 &: \text{SideEffects } 3 \rightarrow 1, \\ p_6 &: \text{SideEffects } 2 \rightarrow 0. \end{aligned}$$

Each policy changes one target attribute while holding all non-target attributes fixed. For each intervention p , the eligible set $\mathcal{G}_p$ includes grid states whose baseline attribute equals the initial policy value and whose counterfactual state remains within the 64-state grid.

I.3 Policy-Level Average Effects

For model m , the average effect of policy p is computed as

$$\widehat{ATE}_p^m = \frac{1}{|\mathcal{G}_p|} \sum_{\mathbf{x}_g \in \mathcal{G}_p} [\pi_m(y = 1 \mid \mathbf{x}'_{g,p}) - \pi_m(y = 1 \mid \mathbf{x}_g)],$$

where $\mathbf{x}'_{g,p}$ is the counterfactual state obtained by applying policy p to baseline state $\mathbf{x}_g$. The evaluated probability functions are

$$\pi_m \in \{\pi_{\text{LLM}}, \pi_{\text{BDT}}, \pi_{\text{DCE}}\}.$$

Static-BDT probabilities are computed using the main anchoring specification:

$$\pi_{\text{BDT}} = 0.25 \bar{\pi}_{\text{LLM}} + 0.75 P_{\text{static}}.$$

I.4 Pairwise Ranking Rule

Let

$$\mathcal{P}_2 = \{(p, q) : p, q \in \mathcal{P}, p < q\}$$

denote the set of unordered policy pairs. With six candidate interventions, $|\mathcal{P}_2| = 15$. For each model m , PRC is computed as

$$\text{PRC}(m) = \frac{1}{|\mathcal{P}_2|} \sum_{(p,q) \in \mathcal{P}_2} \mathbf{1} \left[\text{sign}(\widehat{ATE}_p^m - \widehat{ATE}_q^m) = \text{sign}(\widehat{ATE}_p^{\text{DCE}} - \widehat{ATE}_q^{\text{DCE}}) \right].$$

Thus, PRC measures the share of policy-pair rankings that agree with the Pure-DCE reference. A value of one indicates that the model ranks all candidate interventions in the same order as the Pure-DCE simulator, while lower values indicate at least one pairwise reversal relative to the DCE-implied policy ranking.

I.5 Uncertainty Intervals

No confidence intervals are reported for PRC. The 15 unordered policy-pair comparisons share six underlying policy effects and are not 15 independent Bernoulli trials; the binomial CI and one-sided test framework that would normally apply to a 15-trial success count does not hold here. PRC should be read as a descriptive share of policy pairs with matching rank order to Pure-DCE, not as a sample statistic with inferential uncertainty.

I.6 Reproducibility Checks

The PRC computation is conducted on the same fixed 64-state grid used in the main frontier-alignment analysis. For each model, the analysis verifies that all 64 states are present, that the Pure-DCE probability equals P_{static} , and that the Static-BDT probability satisfies

$$\pi_{\text{BDT}} = 0.25 \bar{\pi}_{\text{LLM}} + 0.75 P_{\text{static}}$$

up to numerical tolerance. The resulting PRC values are reported in Table S10. The metric should be interpreted as a structural diagnostic of the 75% P_{emp} blend, not as external validation against realized vaccine uptake and not as a hypothesis test.

J Bootstrap Uncertainty Intervals

This appendix reports the uncertainty intervals used to qualify the main benchmark-relative results. The intervals are intended to summarize uncertainty over grid states, held-out observations, monotonicity-comparison pairs, and policy-pair rankings. They should not be interpreted as respondent-level sampling uncertainty from re-estimating the full 6-parameter conditional-logit frontier unless explicitly stated.

J.1 Bootstrap Procedure

To quantify uncertainty, we compute bootstrap confidence intervals tailored to the specific unit of analysis for each evaluation metric.

For **held-out prediction metrics** (Table S6 in this supplement, mirroring Table 1 of the main text), the DCE data possess a nested panel structure where multiple choice observations are clustered within individual respondents. To strictly account for within-respondent correlation and avoid

underestimating standard errors, we employ a **respondent-level cluster bootstrap**. Specifically, we resample the 205 unique respondents (whose choices constitute the 813 alternative-level binary-outcome rows) with replacement, maintaining the intact within-respondent choice structure. We perform 2,000 replications to ensure stable interval estimation. Notably, the resulting respondent-level confidence intervals are narrower than their observation-level counterparts. This counterintuitive but statistically sound result occurs because, in our specific data structure (205 clusters with approximately 4 observations each), structured cluster resampling provides a more stable variance estimate. It avoids the extreme duplication noise inherent in observation-level resampling, which can artificially inflate variance by breaking the within-respondent panel structure.

For **frontier-deviation metrics** (Table S11), the evaluation operates on the fixed 64-state policy grid rather than respondent-level choices. Therefore, we compute frontier-deviation intervals by resampling the 64 policy-grid states with replacement using 1,000 replications. For **monotonicity violation rates** (MVR), we resample comparable attribute pairs. For **policy-ranking consistency** (Table S10 in this supplement, discussed in the main text Policy Ranking Consistency section), we report no confidence intervals and no significance test, because the 15 unordered policy pairs share six underlying policy effects and are not 15 independent Bernoulli trials; the metric is reported as a descriptive share only.

These intervals provide uncertainty summaries for the finite evaluation objects used in the paper: policy-grid states, held-out matched respondents, monotonicity-comparison pairs, and policy-pair rankings. They are not a substitute for external validation against realized uptake or for respondent-level uncertainty from re-estimating the full 6-parameter conditional-logit DCE frontier.

J.2 Frontier-Alignment Uncertainty

Table S11 reports bootstrap intervals for the primary frontier-alignment diagnostics discussed in the main text. The intervals support the main Qwen2.5-72B-Instruct alignment result: the MSE and Spearman ρ intervals do not overlap between Static-BDT and the unconstrained LLM. The MVR-Wait point estimates for Qwen2.5-72B-Instruct also do not overlap (raw 0.6250, Static-BDT 0.3021); the bootstrap intervals, however, are wide and overlap on the 64-state grid (state-level resampling leaves the 16 (eff, se) cell pattern close to identical across replicates), so the MVR-Wait improvement is best read from the point estimates and the cross-model ranking rather than from a within-Qwen significance test. The side-effect monotonicity diagnostic is weaker: for MiroThinker-1.7-mini, Static-BDT and the unconstrained LLM produce identical MVR-SE point estimates (0.3333) and identical bootstrap CIs [0.240, 0.427], indicating that the Static-BDT anchor at $\lambda = 0.25$ does not transmit a sufficient side-effect signal on this frontier.

J.3 Held-Out Prediction Uncertainty

Table S6 (this supplement, mirroring Table 1 of the main text) reports bootstrap intervals for the held-out prediction metrics. On the LLM-matched held-out subset, Static-BDT improves substantially over the unconstrained LLM for both log loss and Brier score. The intervals for Static-BDT and the unconstrained LLM do not overlap for either metric. In contrast, the intervals for Static-BDT and Pure-DCE overlap for both log loss and Brier score, supporting the interpretation that Static-BDT brings raw LLM probabilities close to the structured DCE benchmark but should not be interpreted as outperforming Pure-DCE within the DCE attribute space.

K Covariate Balance Check: Matched vs. Unmatched Subset

To assess the representativeness of the 22.1% matched held-out subset ($N = 813$ *alternative-level binary-outcome rows* from $N = 205$ held-out respondents (each row is one of three alternatives within one of six tasks; the held-out pool contains $N = 3,690$ alt-rows; after excluding one anomalous wait=2 alt-row, 3,689 rows remain under the intended design, of which 813 are matched and 2,876 are unmatched)) used in Table 1 of the main text and in Table S6 of this supplement, we compare its covariate distribution against the unmatched held-out alt-row subset ($N = 2,876$). Table S13 reports the means and Standardized Mean Differences (SMD) for key DCE attributes and the choice rate.

The results confirm that the matched subset systematically differs from the unmatched subset on all observed dimensions (all SMDs > 0.1). The root cause is structural rather than sampling-based: the 64-state LLM grid evaluates only specific integer attribute levels, inherently oversampling longer waits, higher efficacy, and higher side-effect burdens. Consequently, the matched subset serves as a harder, non-representative test bed for the specific attribute profiles covered by the grid. Importantly, because all 205 respondents are present in both subsets, this covariate imbalance does not compromise the respondent-level cluster bootstrap intervals reported in the main text.

L Subgroup Heterogeneity Analysis

To assess whether the hybrid framework’s performance is uniform across respondent subgroups, we replicated the held-out predictive validation on the 813 *alternative-level* rows subset (from 205 held-out respondents) stratified by age, gender, and education. Table S14 reports the log loss, Brier score, and Spearman rank correlation of the Static-BDT anchored predictions within each subgroup, alongside the unconstrained Qwen2.5-72B-Instruct baseline for direct comparison. Because the 813 *alternative-level* rows subset is too small for stable within-subgroup bootstrap inference, the table reports point estimates only; bootstrap intervals on the full sample are reported in the main text.

Two patterns emerge. First, the unconstrained LLM produces near-uniform log loss and Brier score across all subgroups (range ≤ 0.05 on both metrics), confirming that the behavioral unreliability documented in the main text is not driven by a particular demographic segment. Second, the Static-BDT Spearman correlation against the within-stratum Pure-DCE predictions falls in a band of 0.897–0.950 across all seven subgroups (range 0.053), with no subgroup falling below 0.89. This near-uniform behavior is consistent with the Static-BDT anchoring rule being a global probability-blending operation that does not condition on respondent characteristics; the empirical frontier itself is population-averaged by construction, and the resulting anchoring benefit is therefore expected to transfer across subgroups defined by age, gender, and education within this sample.

These subgroup results should not be over-interpreted. The 813 *alternative-level* rows subset (from 205 held-out respondents) does not provide sufficient statistical power to detect small subgroup-specific effects, and the Wuhan sample is not designed to be demographically representative of any specific target population. The subgroup analysis is presented as a robustness check on the homogeneity of the anchoring effect, not as evidence for the generalizability of the empirical frontier itself. Future work with larger and more diverse samples is required to assess whether the framework’s anchoring benefit is preserved across demographic groups in other populations.

M Replication Materials

Replication materials are available in a public GitHub repository: <https://github.com/REDACTED/behavioral-digital-twins>. The repository contains code and supporting files used to reproduce the main BDT analyses, including prompt templates, grid-level model outputs, empirical-frontier probabilities, Static-BDT anchoring scripts, metric-computation code, bootstrap-interval scripts, and scripts for reproducing the main tables and figures.

Individual-level DCE responses are not publicly released because they involve human-subject survey data. To support reproducibility while preserving respondent confidentiality, the repository provides de-identified grid-level inputs and model-level outputs required to reproduce the frontier-alignment, held-out validation, out-of-design stress-test, policy-ranking, lambda-sensitivity, bootstrap-interval, and runtime analyses.

N Out-of-Design Narrative Evaluation (Supplementary Table Reference)

The main text (Section 4.3, Table 2 of the main text) reports the out-of-design Panel A and Panel B narrative-policy evaluation on the 20 paired scenarios (40 narrative descriptions) with MVR-Wait on the 8 wait-varying pairs (Panel A: 6; Panel B: 2), MAD/ ρ on the 12 narrative-only pairs, and the two-model comparison (Qwen2.5-72B-Instruct and DeepSeek V4; MiroThinker-1.7-mini was not administered on the OOD panel because its 64-state grid evaluation was used to validate the Static-BDT anchor as a standalone test, not to run the OOD narrative scenarios). The four-panel diagnostic figure (MVR-Wait / Spearman ρ / subgroup heterogeneity / PRC) that combined the 20 paired OOD scenarios with the 813-row matched-subset subgroup analysis is omitted from this supplement because the rank-correlation panel falls inside the construction-only null band; the table-based diagnostics (Table 2 of the main text and Table S14 of this supplement) are the authoritative source for the OOD narrative results. The out-of-design evaluation is restricted to the two-model Qwen2.5-72B-Instruct and DeepSeek V4 panel; the third model (MiroThinker-1.7-mini) is reported only on the 64-state grid in the main text.

O Task-Level Multinomial Held-Out Validation

The 813 alternative-level rows reported in Section 4.2 of the main text use the alt-level binary log loss $\ell(y, p) = -y \log p - (1 - y) \log(1 - p)$, where each row is one of three alternatives within one of six choice tasks per respondent. A natural reviewer question is whether the alt-level binary loss is the right proper score for a three-alternative choice task. The binary score implicitly treats each alternative as an independent binary outcome, whereas the genuine generative model is a three-alternative categorical choice in which $P(A) + P(B) + P(C) = 1$ within each task. We address this directly with a complementary task-level multinomial held-out validation.

O.1 Setup

We hold out 20% of respondents (205 of 1,027, seed = 2026) and refit the corrected 6-parameter conditional logit on the remaining 80% (822 respondents). The three methods being compared differ in how the task-level probability $\Pr(i \mid \text{task})$ is constructed for $i \in \{A, B, C\}$, and we now make the per-alternative and per-task constructions explicit:

Method 1 (Pure-DCE). For each test respondent task, the conditional logit supplies utility U_A, U_B, U_C . The task-level multinomial probabilities are computed by the strict conditional-logit softmax

$$\Pr_{\text{DCE}}(i \mid \text{task}) = \frac{\exp[U_i]}{\sum_{j \in \{A, B, C\}} \exp[U_j]},$$

Taking the opt-out alternative as the utility reference, subtracting its utility from all three alternatives yields utilities $U_A - \hat{\beta}_{\text{ASC}}, U_B - \hat{\beta}_{\text{ASC}}$, and 0; this common normalization leaves softmax probabilities unchanged.

Method 2 (Unconstrained LLM). For each test respondent task, the matched (wait, eff, se) cell supplies a per-alternative LLM acceptance probability $p_{\text{LLM},i} \in [0, 1]$ (averaged over 10 repeated LLM calls in the unconstrained Qwen2.5-72B-Instruct prompting condition). The LLM provides binary acceptance probabilities for A and B (and a placeholder probability for C, which is $\Pr(C) = 0$ for the unconstrained condition because the prompt elicits a yes/no decision against the offered policy). The task-level probability is the renormalized binary blend

$$\Pr_{\text{LLM}}(i \mid \text{task}) = \frac{p_{\text{LLM},i} \cdot \mathbf{1}[i \in \{A, B\}]}{\sum_{j \in \{A, B\}} p_{\text{LLM},j}},$$

which is a softmax over the two non-opt-out alternatives only; this matches how the held-out log-loss code consumes the LLM’s binary outputs.

Method 3 (Static-BDT / EFR). The Static-BDT prediction per alternative is constructed as a per-alternative binary acceptance probability blend

$$p_{\text{BDT},i} = \lambda \cdot p_{\text{LLM},i} + (1 - \lambda) \cdot p_{\text{DCE},i}, \quad \lambda = 0.25,$$

where $p_{\text{DCE},i} = \sigma[U_i - \hat{\beta}_{\text{ASC}}]$ is the same empirical-frontier binary probability used in the alt-level binary validation. Because $p_{\text{BDT},i}$ is a binary acceptance probability rather than a softmax weight, the task-level probability is the renormalized blend

$$\Pr_{\text{BDT}}(i \mid \text{task}) = \frac{p_{\text{BDT},i} \cdot \mathbf{1}[i \in \{A, B\}]}{\sum_{j \in \{A, B\}} p_{\text{BDT},j}},$$

i.e. the convex blend is taken per alternative in the alt-level binary space, then the task-level score is obtained by normalizing the two non-opt-out alternatives to sum to 1. This is the construction that the held-out code implements for both the 813 alt-level rows and the 78 task-level matched tasks; we report it explicitly here so the reviewer does not need to guess.

The LLM lookup is restricted to the 11 (wait, eff, se) cells in the 12-cell candidate superset (wait $\in \{0, 6\}$, eff $\in \{0.5, 0.7\}$, se $\in \{1, 2, 3\}$) that are actually populated in the matched subset (the full 12-cell candidate superset is the literal Cartesian overlap on the three focal attributes between the intended DCE non-optout wait/eff/se support and the 64-state grid; 1 of these 12 cells, namely ($\text{wait} = 0, \text{eff} = 0.7, \text{se} = 2$), is not observed in the matched non-optout pool; 11 cells remain). The held-out pool contains 3,690 alt-rows; after excluding one anomalous wait=2 alt-row, 3,689 rows remain under the intended design, of which 813 are matched and 2,876 are unmatched. Only the 813 matched alt-rows are used for the alt-level binary validation; unmatched tasks are excluded from the LLM and Static-BDT comparison.

O.2 Two-Alternative Subset (A versus B, opt-out excluded)

Restricting to test tasks in which the chosen alternative is A or B (excluding opt-out chosen), and restricting to the LLM-matched subset of the held-out 20% respondent split under the canonical 11-

of-12-cell populated intersection ($n_{\text{tasks, matched}} = 78$), the task-level multinomial log loss evaluated on the *same* $n = 78$ matched tasks for all three methods is:

On this apples-to-apples matched subset, Static-BDT (0.4307) sits *between* the two baselines, reducing task-level log loss by 49.5% relative to the unconstrained LLM (0.8521) but remaining 30.2% above the empirical-anchor Pure-DCE frontier (0.3307): the small blend weight 0.25 is sufficient to absorb the LLM’s worst monotonicity violations, but the LLM component still pulls predictions away from the empirical frontier in the task-level proper score. The two proper scores are not directly comparable in absolute terms—the alt-level binary score treats the unmatched alternatives as missing not at random (i.e., only the matched alternatives within a 2-alternative non-opt-out task contribute a loss term, and the loss is averaged over those rows), while the task-level multinomial score renormalises the two non-opt-out alternatives to sum to 1 and computes the joint log loss over the (A, B) choice. Both proper scores are unbounded above, with the uniform-prediction reference giving a baseline of $\log 2 \approx 0.693$ for the binary case and a comparable baseline for the two-alternative multinomial; we report all three numerical values (binary 0.6789 / 0.6985 / 0.8273 and task-level 0.3307 / 0.4307 / 0.8521) directly without any rescaling, so that the absolute magnitudes can be read against the same uniform-prediction baseline.

O.3 Three-Alternative Full Tasks (A, B, C)

The strict 11-of-12-cell populated intersection has no opt-out alternative in its support, because the opt-out alternative is by construction at $(w, e, s) = (0, 0, 0)$ and the grid cells all have $e \in \{0.5, 0.7\} \neq 0$. Consequently the three-alternative full-task subset, which requires all three alternatives (A, B, C) to lie in the LLM-matched cell set, is empty under the canonical match ($n_{\text{tasks, matched}} = 0$ for the 3-alt full subset; $n_{\text{AB chosen, matched}} = 78$ for the 2-alt subset, which is the same set as above). A coarser match that admits the opt-out placeholder cell by relaxing the efficacy constraint to $e \in \{0, 0.5, 0.7\}$ would recover a non-empty 3-alt full subset, but that relaxation would also admit the opt-out placeholder’s $(0, 0, 0)$ which carries no real attribute information, and so we report the canonical 2-alt subset as the primary apples-to-apples task-level score.

Under the canonical 11-of-12-cell populated intersection the 3-alt full subset is empty, so the apples-to-apples task-level score is the 78-task 2-alt subset from Table S15.

O.4 Implication for the 64-State Validation

The 64-state grid evaluation in the main text uses a single probability per (wait, eff, se) cell, with cash and vaccine-origin attributes held at their reference values (cash = 0, origin = 0). The held-out DCE tasks administer cash $\in \{50, 200, 800\}$ and origin $\in \{0, 1\}$, so the LLM lookup matches 78 of the 1,236 test tasks (6.3%) for the 2-alternative subset under the canonical 11-of-12-cell populated intersection. For these matched tasks, the task-level multinomial validation gives a more nuanced picture than the alt-level binary pattern: Static-BDT ($\lambda = 0.25$) reduces task-level log loss by 49.5% relative to the unconstrained LLM (0.4307 vs. 0.8521) but remains 30.2% above the empirical-anchor Pure-DCE frontier (0.3307). On the alt-level binary 813-row matched subsample (a different aggregation that retains the 2-alts-per-row structure used in the main text), Static-BDT is essentially tied with Pure-DCE (0.6985 vs. 0.6789); the residual 2.9% gap is the same direction as in the task-level score and consistent with a small blend weight. For the unmatched test tasks, the frontier-only Pure-DCE prediction remains the only validated method, and the LLM and Static-BDT entries are not directly comparable because the LLM is not administered at the held-out (wait, eff, se) cell with non-reference cash and origin values.

Table S11: Bootstrap intervals for selected frontier-alignment diagnostics (2000 state-level bootstrap replicates with replacement; MVR-Wait reported as point estimate because the 16 (eff,se) cell-level structure makes the MVR a step function on the grid).

Model	Method	Metric	Estimate	
Qwen2.5-72B-Instruct	Unconstrained LLM	MSE	0.0258	[0.0174, 0.0350]
Qwen2.5-72B-Instruct	Static-BDT Anchor	MSE	0.0016	[0.0011, 0.0022]
Qwen2.5-72B-Instruct	Unconstrained LLM	MVR-Wait	0.6250	point
Qwen2.5-72B-Instruct	Static-BDT Anchor	MVR-Wait	0.3021	point
Qwen2.5-72B-Instruct	Unconstrained LLM	Spearman ρ	0.0608	[−0.1784, 0.3042]
Qwen2.5-72B-Instruct	Static-BDT Anchor	Spearman ρ	0.6886	[0.5266, 0.8012]
DeepSeek V4	Unconstrained LLM	MSE	0.0379	[0.0266, 0.0505]
DeepSeek V4	Static-BDT Anchor	MSE	0.0024	[0.0017, 0.0031]
DeepSeek V4	Unconstrained LLM	MVR-Wait	0.3438	point
DeepSeek V4	Static-BDT Anchor	MVR-Wait	0.0729	point
DeepSeek V4	Unconstrained LLM	Spearman ρ	0.3025	[0.0714, 0.5123]
DeepSeek V4	Static-BDT Anchor	Spearman ρ	0.8123	[0.6832, 0.8885]
MiroThinker-1.7-mini	Unconstrained LLM	MSE	0.0872	[0.0662, 0.1107]
MiroThinker-1.7-mini	Static-BDT Anchor	MSE	0.0054	[0.0041, 0.0070]
MiroThinker-1.7-mini	Unconstrained LLM	MVR-Wait	0.1042	point
MiroThinker-1.7-mini	Static-BDT Anchor	MVR-Wait	0.0833	point
MiroThinker-1.7-mini	Unconstrained LLM	Spearman ρ	0.5306	[0.3070, 0.6966]
MiroThinker-1.7-mini	Static-BDT Anchor	Spearman ρ	0.7564	[0.5849, 0.8860]

Note. Unlike the held-out prediction metrics in Table S6 (this supplement, mirroring Table 1 of the main text), which employ respondent-level cluster bootstrapping to account for the nested panel structure of DCE choices, the frontier-alignment metrics here are evaluated on the fixed 64-state policy grid. Therefore, frontier-deviation intervals are computed by resampling the 64 policy-grid states with replacement (2,000 replications). MVR-Wait is reported as a point estimate: the 16 (eff,se) cell-level structure makes the MVR a step function on the grid, so a state-level bootstrap yields degenerate intervals; a comparable-attribute-pair bootstrap (resampling (eff,se) cells) gives the same point estimate by construction because the frontier-deviation ranking is invariant under cell-level resampling. The MiroThinker-1.7-mini side-effect diagnostic (MVR-SE = 0.3333 for both unconstrained and Static-BDT) is omitted from this table: the corrected 6-parameter specification (Section 3 of the main text) makes the side-effect coefficient the smallest empirical signal among the five attributes, and the Static-BDT anchor at $\lambda = 0.25$ does not transmit a sufficient side-effect correction to move MiroThinker’s MVR-SE off the unconstrained value; the failure mode is documented in Section 4.3 of the main text (OOD panel-level diagnostics note).

Table S12: Attribute-level mapping between the administered DCE design and the 64-state LLM evaluation grid. The matched subset is restricted to cells where the DCE design level and the grid level coincide as a `float` value (no rescaling, no interpolation).

Attribute	DCE Levels	Grid Levels	Literal Intersection (matched subset)
Wait Time (months)	{0, 1, 3, 6}	{0, 2, 4, 6}	{0, 6}
Vaccine Efficacy	{0, 0.5, 0.7, 0.95}	{0.3, 0.5, 0.7, 0.9}	{0.5, 0.7}
Side-Effect Burden (coded level)	{0, 1, 2, 3}	{0, 1, 2, 3}	{1, 2, 3}
Cash Incentives (RMB)	{0, 50, 200, 800}	0 (held)	— (not matched)
Vaccine Origin	{0, 1}	0 (held)	— (not matched)

Table S13: Covariate balance between matched and unmatched held-out subsets.

Variable	Matched Mean	Unmatched Mean	SMD
Wait Time (months)	2.77	1.33	0.58
Vaccine Efficacy	0.59	0.45	0.49
Side Effects	1.95	1.16	0.79
Cash Incentives (RMB)	325	207	0.38
Vaccine Origin (0=dom, 1=imp)	0.50	0.28	0.45
Choice Rate ($y = 1$)	0.528	0.278	0.53

Note. The matched subset consists of held-out observations whose $(wait, eff, se)$ cell is in the literal overlap between the intended DCE design and the LLM grid on the three focal attributes (waiting time, efficacy, and side-effect burden). SMD denotes the Standardized Mean Difference. The DCE design administers $(wait, eff, se) \in \{0, 1, 3, 6\} \times \{0, 0.5, 0.7, 0.95\} \times \{0, 1, 2, 3\}$ (the four real wait design levels; the raw coded data contain one additional 1-row outlier at $wait = 2$ months from a single respondent which is excluded as a data anomaly because it is not an administered design level); the 64-state grid uses $(wait, eff, se) \in \{0, 2, 4, 6\} \times \{0.3, 0.5, 0.7, 0.9\} \times \{0, 1, 2, 3\}$. The literal overlap on the three focal attributes (cash incentives and vaccine origin are held at grid reference values and do not participate in the matching criterion) is $wait \in \{0, 6\}$, $eff \in \{0.5, 0.7\}$, $se \in \{1, 2, 3\}$, giving $2 \times 2 \times 3 = 12$ candidate cells of which 11 are populated in the matched subset (the $(wait = 0, eff = 0.7, se = 2)$ cell is not observed in the matched non-optout pool). Cash incentives (DCE levels $\{0, 50, 200, 800\}$, held at grid reference value 0) and vaccine origin (DCE levels $\{0, 1\}$, held at grid reference value 0) are not part of the matching criterion; DCE cells with $cash > 0$ or $origin = 1$ are matched on $(wait, eff, se)$ alone. The matched subset ($N = 813$ alternative-level rows from 78 task-level 2-alternative non-opt-out tasks) is therefore not a random sample of the unmatched subset ($N = 2,876$ alt-rows under intended design); the SMDs in the table above quantify the systematic differences. Crucially, all 205 held-out respondents contribute to both subsets; thus, the imbalance is strictly at the observation (attribute-level) rather than the respondent level.

Table S14: Subgroup heterogeneity analysis of Static-BDT ($\lambda = 0.25$) performance on the 813 *alternative-level* rows subset (from 205 held-out respondents), stratified by demographic characteristics.

Subgroup	Method	Log loss ↓	Brier ↓	Spearman ρ ↑
Age < 55	Unconstrained LLM	0.835	0.317	−0.628
	Static-BDT	0.698	0.252	0.916
Age ≥ 55	Unconstrained LLM	0.801	0.302	−0.617
	Static-BDT	0.701	0.253	0.922
Female	Unconstrained LLM	0.815	0.308	−0.627
	Static-BDT	0.690	0.248	0.897
Male	Unconstrained LLM	0.844	0.322	−0.622
	Static-BDT	0.710	0.258	0.937
College or University	Unconstrained LLM	0.810	0.306	−0.560
	Static-BDT	0.694	0.250	0.950
Below College	Unconstrained LLM	0.832	0.316	−0.642
	Static-BDT	0.700	0.253	0.906
All (matched subset)	Unconstrained LLM	0.827	0.314	−0.626
	Static-BDT	0.699	0.252	0.917

Note. “Below College” includes middle school or below; “College or University” includes high school, college, or graduate degree. “Age < 55” covers 18–54; “Age ≥ 55” covers 55–65 and 65+. The unconstrained LLM column reports Qwen2.5-72B-Instruct predictions before anchoring. The Spearman ρ column reports rank correlation against the same-method Pure-DCE predictions on the same matched subset (which is the relevant anchor-reference for the LLM’s direction within each stratum), not against the population frontier. Across all seven subgroups, the Static-BDT predictions achieve Spearman ρ between 0.897 and 0.950 (range 0.053), indicating that the empirical anchoring benefit is approximately invariant to demographic subgroup within this Wuhan sample.

Table S15: Task-level multinomial held-out log loss on the 2-alternative subset (A vs. B, opt-out excluded), evaluated on the LLM-matched subset of the held-out 20% respondent split (seed 2026, intended design, exclude wait=2 outlier). All three methods are evaluated on the same $n = 78$ matched tasks. Lower is better. The conditional logit is refit on the 80% training split (six parameters: vaccine origin, wait, efficacy, side effects, cash, opt-out ASC).

Method	Log Loss	N (tasks)	Note
Pure-DCE	0.3307	78	Refit 6-param conditional logit, $\Pr(A) + \Pr(B) = 1 - \Pr(C)$ restricted to $\{A, B\}$
Unconstrained LLM	0.8521	78	Direct LLM $\Pr(A)$ and $\Pr(B) = 1 - \Pr(A)$, restricted to $\{A, B\}$
Static-BDT ($\lambda = 0.25$)	0.4307	78	Convex blend $0.25 \cdot \pi_{\text{LLM}} + 0.75 \cdot \pi_{\text{DCE}}$, restricted to $\{A, B\}$

Table S16: Three-alternative full task subset under the canonical 11-of-12-cell match is empty by construction (the opt-out alt at $(0,0,0)$ has $e = 0$ and never lies in the 12-cell candidate superset). We report the 2-alt subset from Table S15 as the primary task-level score, with the LLM probability for C drawn from the refit 6-param conditional logit for completeness. The 2-alt subset excludes the held-out tasks in which the opt-out was chosen, which are dominated by the unconditional $\Pr(C) \approx 0.20$ from the empirical frontier and would inflate the proper score by a uniform $-\log 0.20 \approx 1.61$ contribution that does not distinguish methods.

Method	Log Loss	N (tasks)	Note
Pure-DCE (3-alt softmax; empty under canonical match)	—	0	No full-3-alt task satisfies the 11-of-12-cell canonical match
Unconstrained LLM (2-alt only)	0.8521	78	$\Pr(A)$ and $\Pr(B) = 1 - \Pr(A)$; same 2-alt subset as Table S15
Static-BDT ($\lambda = 0.25$, 3-alt; empty under canonical match)	—	0	Same as Pure-DCE: 0 full-3-alt tasks under canonical match

References

- [1] John F. P. Bridges, A. Brett Hauber, Deborah Marshall, Andrew Lloyd, Lisa A. Prosser, Dean A. Regier, F. Reed Johnson, and Josephine Mauskopf. Conjoint analysis applications in health—How are studies being designed and reported?: An update: Current practices in ISPOR conjoint analysis applications. *Value in Health*, 16(2):287–298, 2013. doi: 10.1016/j.jval.2012.12.012.
- [2] A. Brett Hauber, Juan Marcos González, Catharina G. M. Groothuis-Oudshoorn, Thomas Prior, Deborah A. Marshall, Charles Cunningham, Maarten J. IJzerman, and John F. P. Bridges. Statistical methods for the analysis of discrete choice experiments: A report of the ISPOR conjoint analysis good research practices task force. *Value in Health*, 19(4):300–315, 2016. doi: 10.1016/j.jval.2016.04.004.
- [3] F. Reed Johnson, Esther Lancsar, Deborah Marshall, Vikram Kilambi, Andreas Muehlbacher, Dean A. Regier, Brian W. Bresnahan, Brett J. Kanninen, and John F. P. Bridges. Constructing experimental designs for discrete-choice experiments: Report of the ISPOR conjoint analysis good research practices task force. *Value in Health*, 16(1):3–13, 2013. doi: 10.1016/j.jval.2012.08.223.
- [4] Deborah A. Marshall, Kenny Deal, Yvonne Bombard, Nicole Leighl, Karen V. MacDonald, Maureen Trudeau, C. Daniel Mullins, Nick Bansback, Paula Mohr, James Doorley, and John E. Brazier. Expanding health technology assessments to include effects on the Health of System Partners: A Concept and Example for Surgical Technology. *Value in Health*, 19(8):e1–e2, 2016. doi: 10.1016/j.jval.2016.08.001.